

COM SCI M151B Week 9

Aidan Jan

November 27, 2024

Memory Consistency

Consistency Problem

- if initially $M[Z] = M[Y] = 0$, $x3 = 100$, $x4 = 200$,

```
Core1: lw x1, Z(x0)           Core2: lw x2, Y(x0)
...
Core1: sw x3, Y(x0)           Core2: sw x4, Z(x0)
```

- What are the potential values for $x1$ and $x2$?
- RAW?
- Correctness?
- LSQ?
 - If the addresses are not the same, why wait?
 - Compiler can reorder, too!
- SW and LW (to different addresses) within one core can be reordered and that will mess up assumptions in OTHER cores.
- The timing between cores might be off! (even without ordering).

Memory Consistency vs. Coherency

- Consistency is about ordering of parallel accesses between *different* addresses.
- Coherency is about ordering of parallel accesses to the *same address*.
- How do we fix this?

The Timing Problem

- Timing between two cores is not the hardware or reordering problem, so it's the programmer's job to fix that using synchronization strategies
- We will discuss this later!

The Reordering Problem

- That is the optimization done by hardware and/or compiler and programmer cannot do anything about this!

Memory Model

- We first need to know what are the possibilities.
- **Memory Model:** Interactions between threads and the shared memory. In other words, *what are the possible values for a load?*

Sequential Consistency

1. Processors see their own loads and stores in program order.
2. Processors see others' loads and stores in program order.
3. All processors see the same global load/store ordering.
4. Sequential Consistency is the simplest scenario!
5. SC makes multicore systems indistinguishable from multi-programmed in-order uniprocessor.
6. However, it is **too slow!**
 - Too many restrictions!

Memory Operation Ordering


- Four types of memory operation orderings:
 - $W \rightarrow R$: write must complete before subsequent read.
 - $R \rightarrow R$: read must complete before subsequent read.
 - $R \rightarrow W$: read must complete before subsequent write.
 - $W \rightarrow W$: write must complete before subsequent write.
 - (This is for *different* addresses, so now RAW hazards!)
- **What can we do?**
 - We really need to be able to re-order memory addresses.
 - * Most lines are not even shared!
 - * Store does not need to even wait.
 - * We have the ability to do so!
 - We need to compromise, and use a "weaker" or "relaxed" model.
 - * Not all loads and stores need to be in-order.
- Sequential consistency maintains all four orderings.
- Relaxed memory consistency models allow *certain* orderings to be violated.

Total Store Order Model

- What do we need?
 - Sequential Consistency
- What do we get?
 - Store Buffer Model

The Total Store Order (TSO) Model is as follows:

- Processor P can read B before its write to A is seen by all processors.
 - Each processor can move its own reads in front of its own writes.
 - Remember LSQs?



- NOTE: Read by other processors **cannot** return new value of A until the write to A is observed by all processors.

How TSO helps?

Let $a = 0$, $flag = 0$.

thread 1	thread 2
<pre>store 1 → a store 1 → flag</pre>	<pre>loop: if (flag == 0) goto loop load a</pre>

- Can load read 0?

TSO vs. SC

- if initially, $M[Z] = M[Y] = 0$

Core1: lw x1, Z	Core2: lw x2, Y
...	...
Core1: sw Y, 100	Core2: sw Z, 200

- What are the potential values for x1 and x2?

TSO Summary

- Improving the per-core performance by allowing the reordering within the core!
- Still not observable by other cores (timing problem)

Weak Consistency

- Different assumptions
 - What and when a memory access is visible to the other cores.
- Performance vs. Complexity
- Different architectures use different models.
- More about it soon!

RISC-V Memory Model

- RISC-V OoO (BOOM) follows the RVWMO memory consistency model (RISC-V weak memory ordering).
- Loads are optimistically fired to memory on arrival to the LSU (for performance)
- Simultaneously, the load instruction compares its address with all of the store addresses that it depends on
 - If there is a match, the memory request is killed.
 - If the corresponding store data is present, then the store data is forwarded to the load and the load marks itself as having succeeded.
 - If the store data is not present, then the load goes to sleep.
 - * Loads that have been put to sleep are retried at a later time.
- BOOM currently exhibits the following behavior:
 - Write \rightarrow Read constraint is relaxed (newer loads may execute before older stores).
 - Read \rightarrow Read constraint is maintained (loads to the same address appear in order).
 - A thread can read its own writes early.

How to write a correct program?

- The memory model is given.

The Simple Solution

- Let it be!
 - Race-free programming

The Not-Very-Simple Solution

- Have the hardware to track potential violations!
 - Each core reorders but if another core needs that data, it should recover.
 - Very costly!


More Realistic Solution

- Use synchronization!
 - The programmer (and/or hardware) use some primitives to synchronize different programs.
 - This solves both timing and ordering problems.


Synchronization

- There are two major solutions:
 1. Producer-Consumer Model
 2. Mutual Exclusiveness.

1- Producer-consumer model



2- Mutual exclusiveness




• Solution:

- Use *fences*!
 - * FENCE is an instruction that enforces all memory operations *before* it to complete before fence is executed
 - * All memory operations *after* FENCE must wait for FENCE to be finished.
 - * This happens *within* each core!

FENCE in RISC-V

- FENCE (rwio)
 - fence r, r



What about between cores?

- Use semaphores!

Semaphores

- Create some barriers!
- This will synchronize all the processors.
- Typically implemented in software.

```
Core1: lw x1, Z      Core2: lw x3, Y
...
Core1: sw Y, 100     Core2: sw Z, 200
```

How to implement barrier/semaphore in hardware?

- Hardware-or with one-hot encoding → set a flag to one when a core reaches the barrier. Wait until all flags are one!

Mutual Exclusiveness

- Making sure only one process has access to the shared region at any given time.
- To do this, use Mutex/Locks!

```
Acquire (lock)      Acquire (lock)
Core1: lw x1, Z      Core2: lw x3, Y
...
Core1: sw Y, 100     Core2: sw Z, 200
Release (lock)
```

Release Consistency

- It is guaranteed that all reads and writes are executed and visible to all cores when the lock is released.
- Other than this guarantee, any ordering outside the acquire/release blocks is possible.

How to implement acquire/release?

In code:

```
void acquire(lock) {
    while (lock != 0) {} //spin lock!
    // lock is acquired at the point this executes.
    lock = 1; // set to 1 to make sure no one else can get it
}

void release(lock) {
    lock = 0;
}
```

In assembly:

```

acquire:
    lw ra, 0(x1)
    bnez ra, acquire    # <- loops to the beginning
    addi ra, x0, 1
    sw ra 0(x1)

release:
    st r0 -> 0(&lock)

```

Problem?

<p>(Core 1)</p> <pre> acquire: lw ra, 0(x1) bnez ra, acquire addi ra, x0, 1 sw ra 0(x1) </pre>	<p>(Core 2)</p> <pre> acquire: lw ra, 0(x1) bnez ra, acquire addi ra, x0, 1 sw ra 0(x1) </pre>
--	--

The locks aren't always exclusive! Since the `addi ra, x0, 1` is executed after the check for the lock, two cores can both claim the lock at the same time!

The solution: atomic instructions!

Atomic Instructions

- A hardware primitive that guarantees instructions are executed in-order and without interruption.

How to implement Mutex?

- Hardware-enforce instruction (called test and set).

```

TS(int x) {
    oldval = SWAP(x, 1); // atomic swap
    return oldval;
}

```

- Other variants: compare-and-swap, fetch-and-add

How SWAP is implemented?

- Load-reserved, store-conditional (in RISC-V)

```

Label:
    load-link x2, 0(x1)
    addi x2, x0, 1
    store-conditional x2, 0(x0)
    branch-not-zero label    # check for failure

```

LR and SC

- On load-link, processor remembers address.
- Then checks for writes by other processors.
 - If write is detected, store-conditional will fail.

- For MESI:
 - Load-Reserved ensures line in cache in Exclusive/Modified state.
 - Store-Conditional succeeds if line still in Exclusive/Modified state.

Recap

- Locks:
 - Test-and-set
 - Atomic
 - LR and SC
- This can be combined with a very weak consistency model (called release consistency).

Optimizations?

- Too many SWs when waiting/spinning.
 - Use test-test-and-set
- Unfairness?
 - Use queues
- Multiple/Parallel Locks

Bottom Line:

- There is a range of possibilities!
- The burden is (mostly) on the developer!

Test-test-and-set

```
TTS(int x) {
    oldval = x;
    if (oldval == 0)
        oldval = SWAP(x, 1);
    return oldval;
}
```

- To check for lock:
- `while (TTS(lock) == 1);`